


E3 robustness: over-optimization is smooth in bias and mitigable by the KL penalty

(a) loss grows smoothly with bias



(b) KL leash graded-ly mitigates hacking

