

E5: reward-model over-optimization in a synthetic RLHF-style pipeline (learned RM from biased preferences + PPO + gold eval)

