


E5 sensitivity: reward-model bias-cloning and over-optimization both scale with the systematic bias in the preference signal

(a) the reward model clones more bias as the preference bias grows



(b) over-optimization scales with the preference bias (bias in $\rightarrow$ bias out)

